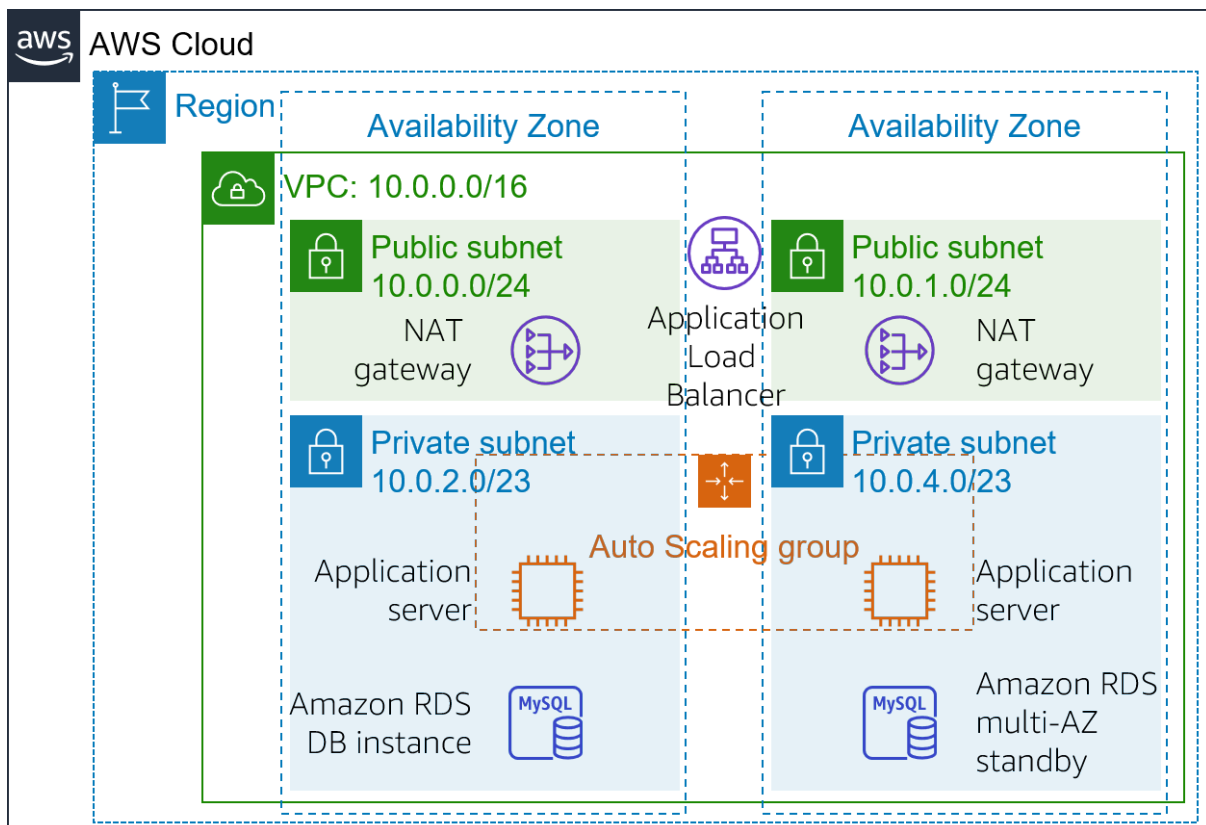


Highly Available Environment using Elastic Load Balancers & Autoscaling Groups

Type	AWS ARCHITECTURE
Tasks Achieved	- Creating and Application Load Balancer - Using an Autoscaling Group to allow scaling of EC2 Instances incase of a failure or in need of more resources

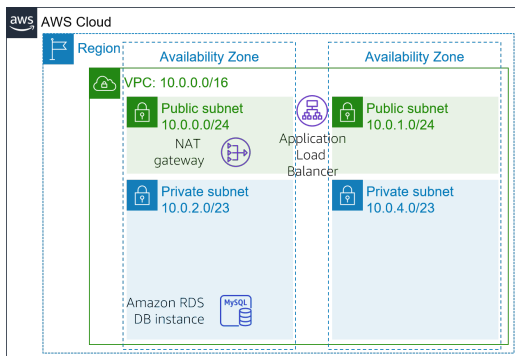
Tasks Achieved



▼ Creating Application Load Balancer - launched in the public subnets

▼ Creating a highly available DB

- Ensure you have a DB subnet group across two availability



→ You will launch the ALB in the public subnets to as they are internet facing

→ The security group will only allow HTTP & HTTPS traffic

→ When setting the listeners you can create or choose from a target group. The target group will be instances which will have port 80 open from which the listeners from the Application Load Balancer will listen from. You can set the following in advance settings:

Health check port
The port the load balancer uses when performing health checks on targets. By default, the health check port is the same as the target group's traffic port. However, you can specify a different port or an override.

☒ Traffic port
☐ Override

Healthy threshold
The number of consecutive health checks successes required before considering an unhealthy target healthy.

5

2-10

Unhealthy threshold
The number of consecutive health check failures required before considering a target unhealthy.

5

2-10

Timeout
The amount of time, in seconds, during which no response means a failed health check.

5 seconds

2-120

Interval
The approximate amount of time between health checks of an individual target.

30 seconds

2-300

Success codes
The HTTP codes to use when checking for a successful response from a target. You can specify multiple values (for example, "200,202") or a range of values (for example, "200-299").

200

→ When setting up the load balancer you choose the two availability zones to launch from as well as the subnet

You will require:

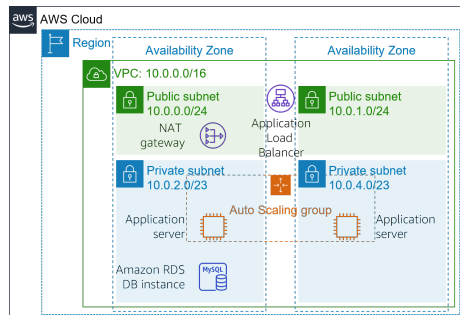
zones by creating a standby instance

- Edit already created DB to allow Multi-AZ deployment
- Remember as you make it highly available double the instances as well as the storage capacity

▼ Creating a highly available NAT Gateway in the case where the one in AZ 1 fails

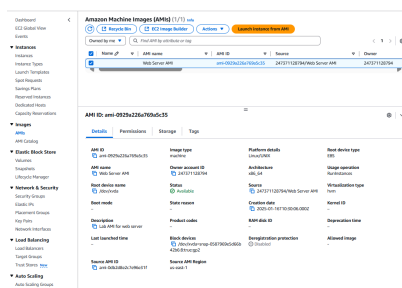
- Remember to allocate an elastic IP when creating
- Create a separate route table to associate the second private subnet to the created NAT gateway

- An autoscaling group that would automatically launch instances in the private subnets

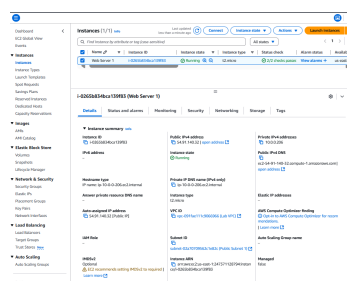


1. Create an AMI from an existing instance

▼ Create an AMI




- An Image ID is a copy of another EC2 Instance. In this case Web Server 1



2. Create a launch template


▼ Choose the AMI from which the instance will be created



- You can add additional setting in the advanced settings section such as the IAM role and the user data

3. An autoscaling group - launched in the private subnets

▼ After creating the launch template you can create an autoscaling group




- **Choose instance launch options:**
Choose the VPC network environment that your instances are launched into, and customize the instance types and purchase options.
 - remember that the instances are being launched in

the private
subnets

- **Integrate with other services** - Use a load balancer to distribute network traffic across multiple servers. Enable service-to-service communications with VPC Lattice. Shift resources away from impaired Availability Zones with zonal shift. You can also customize health check replacements and monitoring.
 - this is where you choose the already created application load balancer
- **Configure group size and scaling** - Define your group's desired capacity and scaling limits. You can optionally add automatic scaling to adjust the size of your group.
- You can choose to add notifications via creating a SNS topic

- Remember to edit the application server inbound rule to allow HTTP/HTTPS traffic from the load balancer as well as for the DB instance to allow DB traffic from the App server



- User data for an inventory website template

```

#!/bin/bash
# Install Apache Web Server and PHP
yum install -y httpd mysql
amazon-linux-extras install -y php7.2
# Download Lab files
wget https://aws-tc-large-objects.s3-us-west-2.amazonaws.com/ILT-TF-200-ACACAD-20-EN/mod9-guided/scripts/inventory-app.zip
unzip inventory-app.zip -d /var/www/html/
# Download and install the AWS SDK for PHP
wget https://github.com/aws/aws-sdk-php/releases/download/3.62.3/aws.zip
unzip aws -d /var/www/html
# Turn on web server
  
```

```
chkconfig httpd on  
service httpd start
```